

RISC-V ARCHITECTURE TRAINING

@DEMO: Bare-metal assembly & SPIKE simulator

Jim Wang (<http://phdbreak99.github.io>)

Dec. 2019

Table of Content

Software stack

GNU toolchain

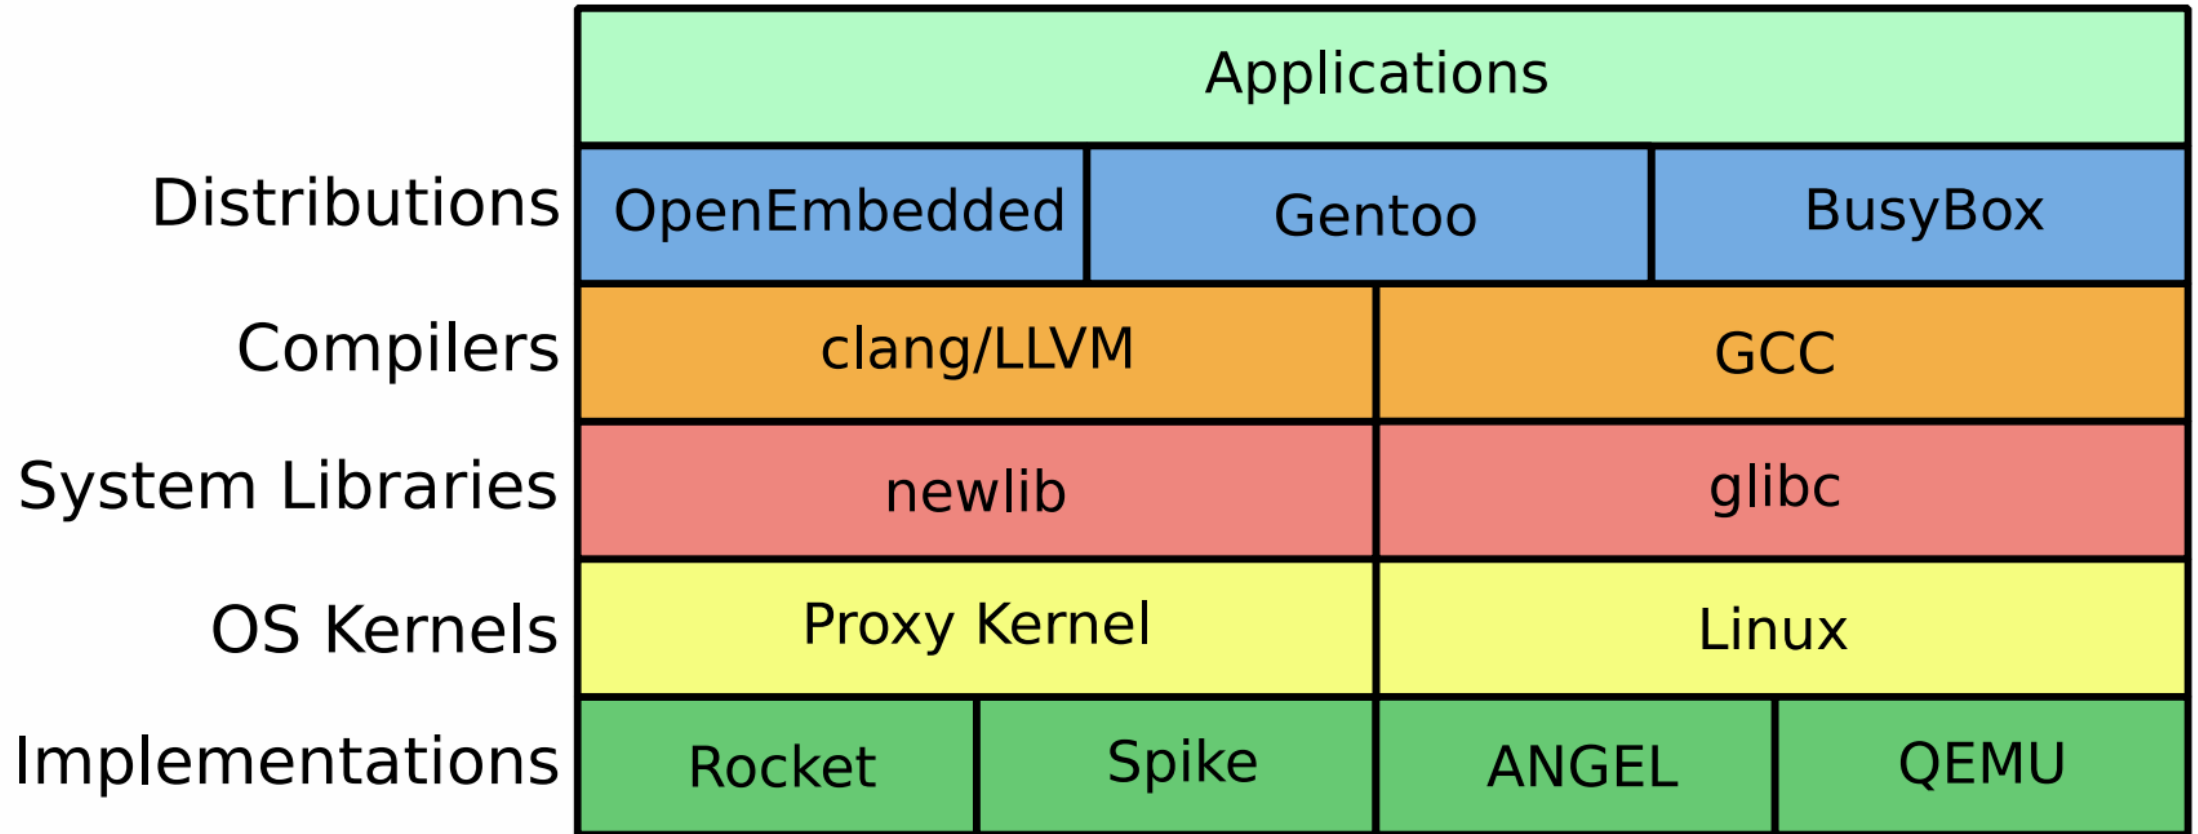
Assembly

SPIKE

Table of Content

>>>> **Software stack**
GNU toolchain
Assembly
SPIKE

General software stack



Embedded system software stack

What is newlib?

<https://en.wikipedia.org/wiki/Newlib>

- C standard library implementation for **embedded system**
- GCC port for non-Linux embedded system
- When lacking of full-blown OS, how to **make a system call** and how to **use devices**

Newlib code size will significant larger than Linux code size, because it includes the system calls that is already embedded inside Linux.

What is cross-compile?

Cross-compiler

- A compiler capable of creating executable code for a platform other than the one on which the compiler is running
- In our case: RISC-V compiler running on top of x86

Table of Content

Software stack

>>>> GNU toolchain

Assembly

SPIKE

Setup GNU toolchain for RISC-V

2 options

1. Build from scratch
2. Download pre-built version from SiFive (or other vendors)

Here we choose option 1, because it's more useful in the future. You probably need to choose your own instruction subsets.

Setup GNU toolchain for RISC-V (cont'd)

1. Download source

- Community version GNU toolchain on Github: <https://github.com/riscv/riscv-gnu-toolchain>
 - riscv-gcc
 - riscv-gdb
 - riscv-glibc
 - riscv-binutil
 - riscv-newlib
 - riscv-dejagnu

```
git clone https://github.com/riscv/riscv-gnu-toolchain --recursive
```


2. Install prerequisites

Ubuntu 16.04

```
sudo apt-get install -y autoconf automake autotools-dev curl libmpc-dev libmpfr-dev libgmp-dev gawk build-e
sudo apt-get install -y build-essential zlib1g-dev pkg-config libglib2.0-dev binutils-dev libboost-all-dev
sudo apt-get install -y autoconf automake autotools-dev curl libmpc-dev libmpfr-dev libgmp-dev libusb-1.0-0
```

Setup GNU toolchain for RISC-V (cont'd)

3. Compile & install

```
git clone --recursive https://github.com/riscv/riscv-gnu-toolchain
# this will take a long time to download

cd riscv-gnu-toolchain; mkdir build; cd build

../configure --prefix=/opt/riscv --with-arch=rv64gc --with-abi=ilp64d
# --with-arch=rv64gc defines target architecture is rv64gc (64-bit IMACFD extensions support)
# option example: rv64imac (64-bit IMAC extensions support)
# --with-abi=ilp64d defines target ABI (application binary interface)
# "d" means hard-float
# option example: ilp64 (64-bit soft float)

make newlib -j4 # compile & install
make report-newlib # run DejaGnu test suite (super slow)
```

Setup GNU toolchain for RISC-V (cont'd)

Toolchain directory content

```
root@8559ed0a43df:~# ls $RISCV/bin
```

elf2hex	riscv64-unknown-elf-cpp	riscv64-unknown-elf-ld.bfd
ivshmem-client	riscv64-unknown-elf-elfedit	riscv64-unknown-elf-nm
ivshmem-server	riscv64-unknown-elf-g++	riscv64-unknown-elf-objcopy
openocd	riscv64-unknown-elf-gcc	riscv64-unknown-elf-objdump
qemu-ga	riscv64-unknown-elf-gcc-8.3.0	riscv64-unknown-elf-ranlib
qemu-img	riscv64-unknown-elf-gcc-ar	riscv64-unknown-elf-readelf
qemu-io	riscv64-unknown-elf-gcc-nm	riscv64-unknown-elf-run
qemu-nbd	riscv64-unknown-elf-gcc-ranlib	riscv64-unknown-elf-size
qemu-riscv32	riscv64-unknown-elf-gcov	riscv64-unknown-elf-strings
qemu-riscv64	riscv64-unknown-elf-gcov-dump	riscv64-unknown-elf-strip
riscv64-unknown-elf-addr2line	riscv64-unknown-elf-gcov-tool	runtest
riscv64-unknown-elf-ar	riscv64-unknown-elf-gdb	spike
riscv64-unknown-elf-as	riscv64-unknown-elf-gdb-add-index	spike-dasm
riscv64-unknown-elf-c++	riscv64-unknown-elf-gprof	termios-xspike
riscv64-unknown-elf-c++filt	riscv64-unknown-elf-ld	xspike

Table of Content

Software stack

GNU toolchain

>>>> Assembly

SPIKE

Assembly / programmer's handbook

Please refer to handouts: *RISC-V Reference Card*

Register	ABI Name	Saver	Description
x0	zero		Hard-wired zero
x1	ra	Caller	Return address
x2	sp	Callee	Stack pointer
x3	gp		Global pointer
x4	tp		Thread pointer
x5-7 & x28-31	t0-6	Caller	Temporaries
x8-9 & x18-27	s0-11	Callee	Saved registers
x10-17	a0-7	Caller	Function arguments / return value

Separation of saved registers and temporary registers makes it possible to reduce 32 registers to 16 registers in E extension

Assembly / what is ABI?

ABI (application binary interface) includes:

- Instruction set
- Calling convention
 - Function's argument passing and return value retrieving
 - Stack vs. registers
 - If stack, which parameter is pushed first?
 - If register, which registers are used for what?
- How to make system calls to operating system
 - More details in our next DEMO

Assembly / `ra` return address

- `ecall`: `ra <= PC + 4`
- `ret`: jump back to `ra` (`PC <= ra`)

Assembly / **sp** stack pointer

When goes into function call, save registers to stack

000000000000114da <_realloc_r>:

114da:	715d	addi	sp,sp,-80	# reserve 80-byte space on stack
114dc:	f84a	sd	s2,48(sp)	# push s2
114de:	e486	sd	ra,72(sp)	# push ra
114e0:	e0a2	sd	s0,64(sp)	# push s0
114e2:	fc26	sd	s1,56(sp)	# push s1
114e4:	f44e	sd	s3,40(sp)	# push s3
...				# push s4 ~ s7
114ee:	e062	sd	s8,0(sp)	# push s8
...				# function
115e0:	60a6	ld	ra,72(sp)	# pop ra
115e2:	6406	ld	s0,64(sp)	# pop s0
115e4:	854a	mv	a0,s2	
115e6:	74e2	ld	s1,56(sp)	# pop s1
...				# pop s2 ~ s7
115f4:	6c02	ld	s8,0(sp)	# pop s8
115f6:	6161	addi	sp,sp,80	# release 80-byte space on stack
115f8:	8082	ret		# return

Assembly / **gp** global pointer

gp = global pointer = pointer to global variables

- GP is pointing at the center of **.data** section that allows program to index to any global variables easily without the need to **auipc** every time

Example: C program uses global variables

```
/* Global Variables: */
Boolean      Bool_Glob;
char         Ch_1_Glob,
            Ch_2_Glob;
Proc_4 () /* without parameters */ {
    Boolean Bool_Loc;
    Bool_Loc = Ch_1_Glob == 'A';
    Bool_Glob = Bool_Loc | Bool_Glob;
    Ch_2_Glob = 'B';
} /* Proc_4 */
```

Assembly / **gp** global pointer (cont'd)

ASM disabled GP

0000000040400826 <Proc_4>:

40400826:	3fc00797	auipc	a5,0x3fc00
4040082a:	f777c783	lbu	a5,-137(a5) # 8000079d <Ch_1_Glob>
4040082e:	3fc00717	auipc	a4,0x3fc00
40400832:	f7272703	lw	a4,-142(a4) # 800007a0 <Bool_Glob>
40400836:	fbf78793	addi	a5,a5,-65
4040083a:	0017b793	seqz	a5,a5
4040083e:	8fd9	or	a5,a5,a4
40400840:	3fc00717	auipc	a4,0x3fc00
40400844:	f6f72023	sw	a5,-160(a4) # 800007a0 <Bool_Glob>
40400848:	3fc00797	auipc	a5,0x3fc00
4040084c:	04200713	li	a4,66
40400850:	f4e78a23	sb	a4,-172(a5) # 8000079c <Ch_2_Glob>
40400854:	8082	ret	

Assembly / **gp** global pointer (cont'd)

ASM enabled GP (cont'd)

```
00000000400003f0 <Proc_4>:
  400003f0: 8651c783          lbu    a5,-1947(gp) # 80001fbd <Ch_1_Glob>
  400003f4: 8681a703          lw     a4,-1944(gp) # 80001fc0 <Bool_Glob>
  400003f8: fbf78793          addi   a5,a5,-65
  400003fc: 0017b793          seqz   a5,a5
  40000400: 00e7e7b3          or     a5,a5,a4
  40000404: 86f1a423          sw     a5,-1944(gp) # 80001fc0 <Bool_Glob>
  40000408: 04200713          li     a4,66
  4000040c: 86e18223          sb     a4,-1948(gp) # 80001fbc <Ch_2_Glob>
  40000410: 00008067          ret
```

Assembly / `tp` thread pointer

`tp` (thread pointer) is a pointer to thread-level global variables (aka thread-local storage)

Assembly / code example

@DEMO

- Directory `~/riscv-training/lab/21-lab.compile`
 - Source code `example-asm.s` and `example-c.s`

Function of `example-asm.s`

- 4x4 Matrix multiplication, and result checking against Excel
- Use 2-level function calls to do the job
 - Demonstrate calling convention by passing argument and return value via registers `a*`
 - Save registers `s*` to stack before using them

Compare with `example-c.c` with the same functionality

- Assembly code is much harder to write and debug for normal functionality
- Assembly code's binary size is smaller (6624 bytes vs. 6000 bytes)

Assembly / what is linker script?

- Describe how the sections in the input files should be mapped into the output file
- Control the memory layout of the output file

Entry point

- The first instruction to execute in the program

Common section

- `.text`: actual machine instructions
- `.data`: static data in your code
- `.bss`: uninitialized global or static variables, will be initialized to zero during startup
 - `.noinit`: part of bss but will not be initialized to zero

Assembly / compile assembly

Compile -> link -> objdump

```
# assemble
${RISCV}/bin/riscv64-unknown-elf-as example-asm.s -o example-asm.o
# link
${RISCV}/bin/riscv64-unknown-elf-ld -T linker-asm.ld example-asm.o -o example-asm.elf
# object dump
${RISCV}/bin/riscv64-unknown-elf-objdump -D example-asm.elf > example-asm.elf.dump
```

Linker script

```
SECTIONS
{
    . = 0x10000;
    .text : { *(.text) }
    .data : { *(.data) }
}

ENTRY (_start)
```

- Both code and data start from `0x0001_0000`
- `_start` is the entry point label

Assembly / compile C code

Compile bare-metal C program

```
# compile  
${RISCV}/bin/riscv64-unknown-elf-gcc example-c.c -o example-c.elf  
# object dump  
${RISCV}/bin/riscv64-unknown-elf-objdump -D example-c.elf > example-c.elf.dump
```


Assembly / ASM vs. C

Development effort

Myself

- 2 hours in ASM
- 2 mins in C

Size of the code

With `printf`

```
riscv@riscv:~/riscv-training/lab/21-lab.compile$ ll *.elf
-rwxr-xr-x 1 1380539737 1876110778 6000 Dec 7 17:07 example-asm.elf*
-rwxr-xr-x 1 1380539737 1876110778 138792 Dec 7 17:18 example-c.elf*
```

Without `printf` and turn on `-Os`: 107.2%

```
-rwxr-xr-x 1 1380539737 1876110778 6000 Dec 7 17:30 example-asm.elf*
-rwxr-xr-x 1 1380539737 1876110778 6432 Dec 7 17:30 example-c.elf*
```

Table of Content

Software stack

GNU toolchain

Assembly

>>>> SPIKE

SPIKE

- SPIKE: official ISS (instruction set simulator) of RISC-V
 - GDB-like TUI (text-based user interface)
 - Support single step execution / breakpoint / watchpoint
 - XSPIKE: open a separate terminal (in GUI mode) to capture the `printf` output

How to invoke SPIKE

```
# run SPIKE in direct mode
> ${RISCV}/bin/spike target.elf

# run SPIKE in interactive debug mode: -d
> ${RISCV}/bin/spike -d target.elf

# run SPIKE with log dumping: -l
> ${RISCV}/bin/spike -l target.elf 2>&1 | less
```

SPIKE (cont'd)

SPIKE interactive debug mode

- `: pc 0`: show current PC in core 0
- `: reg 0 a0`: show content of register `a0` in core 0
- `: mem 2020`: show content of memory at `0x2020`
- `: until pc 0 80000000`: stop when PC hits `0x8000_0000`

More commands type `help` under interactive debug mode

■ Note: don't forget the "0" for core 0

SPIKE (cont'd)

@DEMO

- Run SPIKE in direct mode
- Run SPIKE in interactive debug mode
 - Show register/memory content
 - Set breakpoint
- Run SPIKE with log dumping

감사합니다 Natick
Danke Ευχαριστίες Dalu
Grazie Thank You Köszönöm
Tack
Спасибо Dank Gracias
谢谢 Merci Seé
ありがとう

Obrigado

Next session: LAB

@LAB: factorial in assembly

Use assembly to implement factorial function

$$n! = n * (n-1) * (n-2) * \dots * 2 * 1$$