

RISC-V Training

Uncore

Jim Wang (<http://phdbreak99.github.io>)

July 2019

Table of Content

Uncore

Interrupt Controller

TileLink

Debug

Summary

Table of Content

>>>> Uncore

Interrupt Controller

TileLink

Debug

Summary

Uncore

CPU core is fun, but uncore is the real work.

Uncore / components

- Cache (already discussed)
- Interrupt controller
- Network Fabric
- Debug

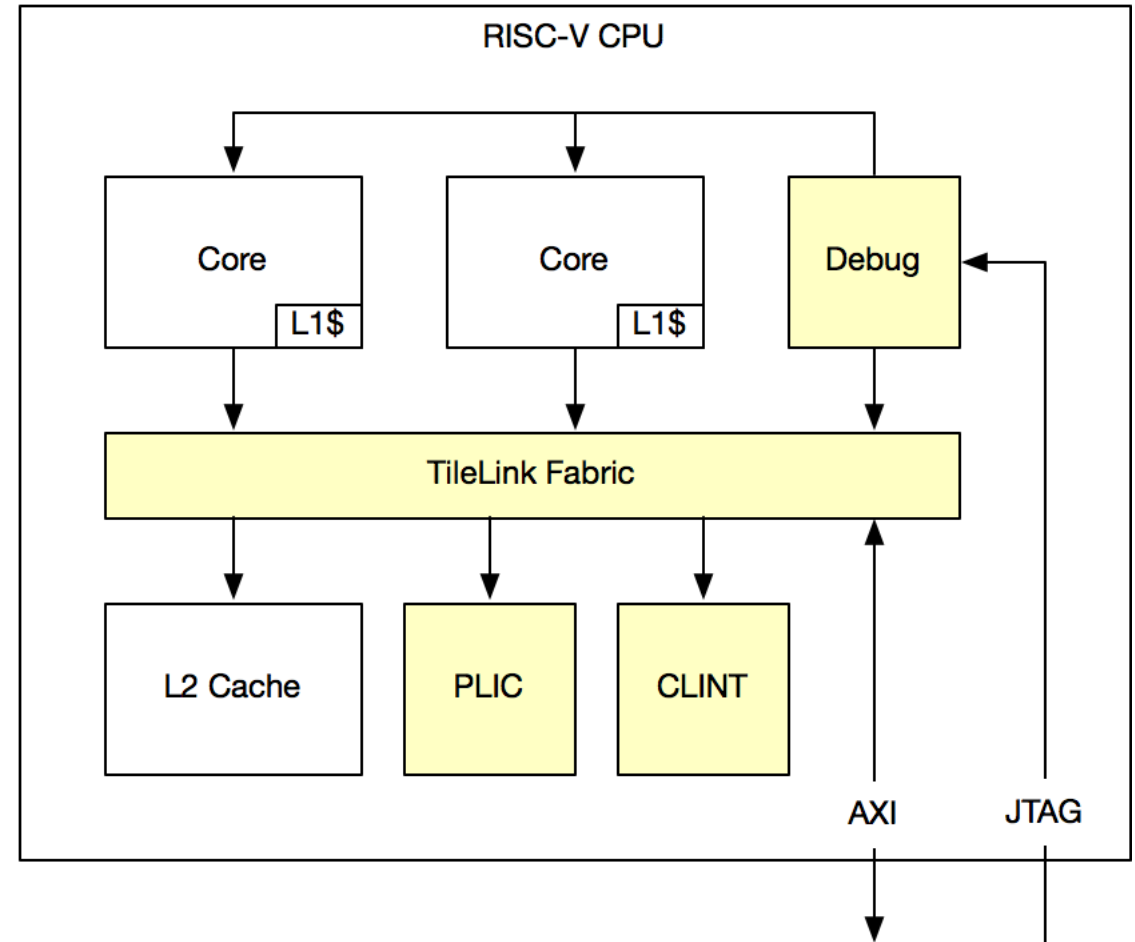


Table of Content

Uncore

>>>> Interrupt Controller

TileLink

Debug

Summary

Interrupt recap

3 types of interrupts

- External: peripheral devices
- Software: inter-processor interrupt
- Timer: timely schedule tasks

PLIC (platform level interrupt controller)

- For external interrupts
- Aggregation of multiple external interrupts
 - Provide enable/priority

CLINT (core level interruptor)

- For software & timer interrupts
- Provide memory-mapped software/timer interrupt CSRs

PLIC (platform-level interrupt controller)

MSI vs. IRQ

MSI (message signaled interrupt)	IRQ (physical wired interrupt)
Relatively newer	Traditional and easy to understand
More complex device and interrupt controller	Simple to implement. Widely supported
Scalable, especially for modern large SoCs	Don't scale well
In-band access, easy for timing/clock/etc.	Out-of-band, nightmare for physical design if many

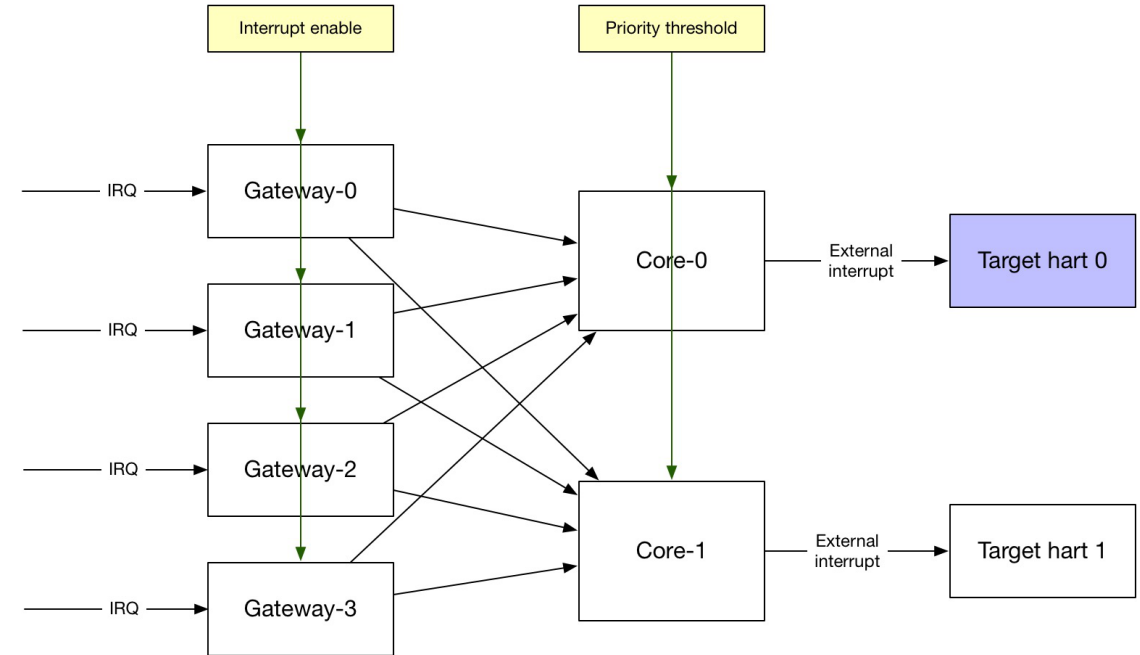
Interrupt controller for RISC-V: PLIC

- Simple & easy to use
- IRQ-style interrupt aggregator
- Support programmable priority and enable
- Support multiple source and multiple target

PLIC / function diagram

IRQ-style interrupt aggregator that support **enable/priority** and **claim/complete** flow

- Function components
 - Gateway
 - Forward enabled interrupts
 - Stop further interrupts if current one is not claimed
 - PLIC core
 - Order interrupts by their priority
 - Filter interrupt with priority threshold

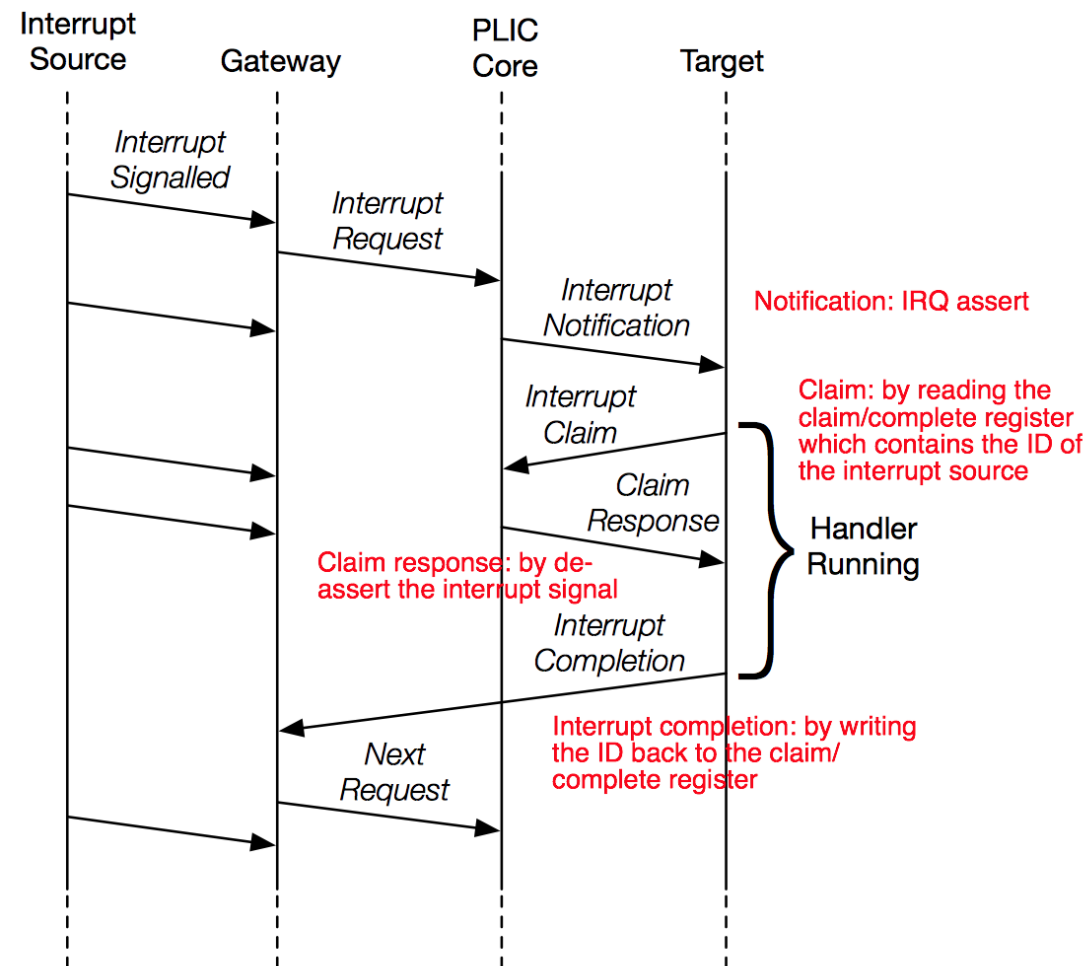


PLIC / configurations

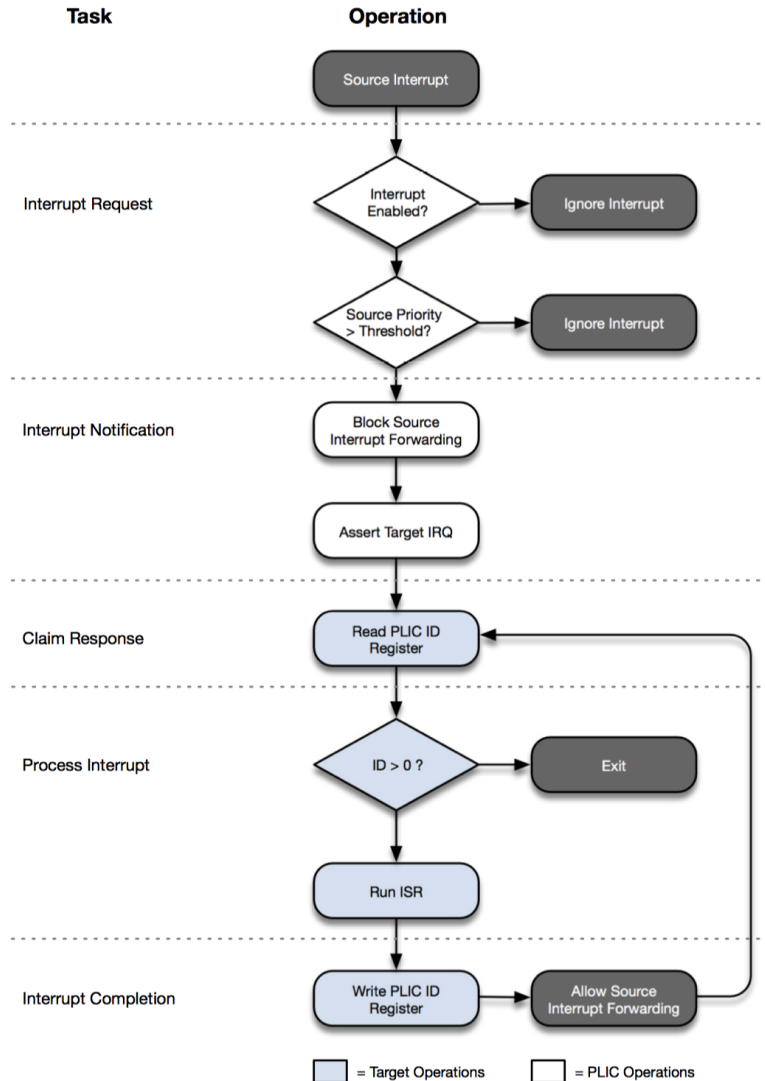
- Interrupt enable
 - Every interrupt initiator has an enable bit
- Interrupt priority
 - Every interrupt initiator has a priority level assigned. By default 8 levels, the larger the number the higher priority it has
 - Level 0 is reserved for “never interrupt”
- Priority threshold
 - Only interrupts have higher priorities will presented to target

PLIC / software-hardware workflow

1. Signalled: IRQ from source to gateway
2. Request: IRQ from gateway to core
3. Notification: IRQ assert from core to target
4. Claim: memory read from target to core, start trap handler
5. Response: IRQ deassert from core to target
6. ... (waiting for trap handler to finish)
7. Complete: memory write from target to core



PLIC / interrupt handler



CLINT

Timer interrupt

- Memory-mapped CSR `mtime` and `mtimecmp`

Software interrupt

- Memory-mapped CSR `msip` (machine-mode software interrupt pending)

CLINT / interrupt handler

Timer interrupt

- Set timer: CPU read `mtime`, then write `mtimecmp` with number larger than `mtime`
- Claim: by writing `mtimecmp` with number smaller than `mtime` to disable/clear timer interrupt

Software interrupt

- Raise software interrupt: write memory-mapped CSR `msip`
- Claim: clear `msip`

Table of Content

Uncore

Interrupt Controller

>>>> TileLink

Debug

Summary

TileLink

Goal of TileLink

- Non-ARM interconnect protocol
 - Cannot use ARM's open protocol: AMBA (CHI, ACE, AXI, AHB, APB)
 - They are open, but fully controlled by ARM. Huge problem for ARM's rival.
- Protocol framework to unify different usage scenarios
 - Clean slate, start from scratch, learn from priors
 - Decouple message protocol from wire protocol
 - Even support off-chip connection
- Simple & hardware-friendly

TileLink is also developed inside UCB BAR.

TileLink is not **tied** with RISC-V. RocketChip uses TileLink.

TileLink / overview

Different usage scenarios

- Point-to-point simple protocol for data transfer
- More complicated data transfer with burst, atomic support
- Cache-coherent hardware support

5 levels of priorities/channels: A, B, C, D, E

- Avoid deadlock: acknowledges have higher priority
- Support out-of-order design for better performance, while support ordering when needed
- Not always need 5 levels of priorities. Simple data transfer will only need 2

Point-to-point, master-slave model

- No race for ownership physically

TileLink / conformance levels

Corresponding to 3 usage scenarios

TL-UL (TileLink uncached light-weight)

- Just read and write operations

TL-UH (TileLink uncached heavy-weight)

- Add burst, atomic support
- Add hint operations, such as prefetch

TL-C (TileLink cached)

- Add support to cache-coherent

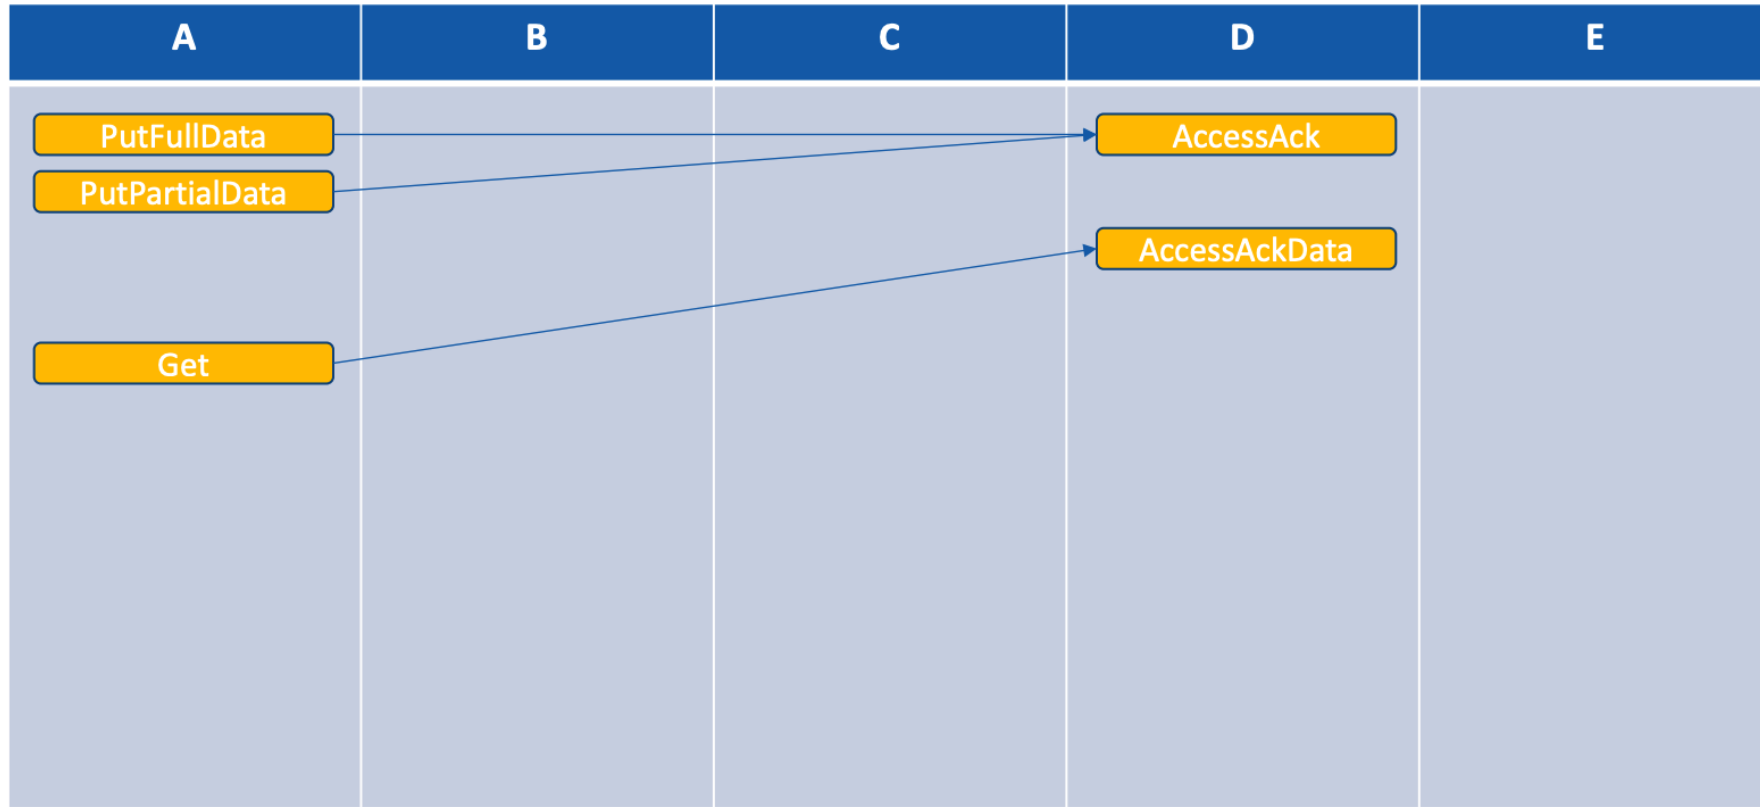
TileLink / channel priorities

Channel A, B, C, D, E

- Each channel has its specified priority
 - In the order of $A < B < C < D < E$
- Channel has only one direction
- Physically independent

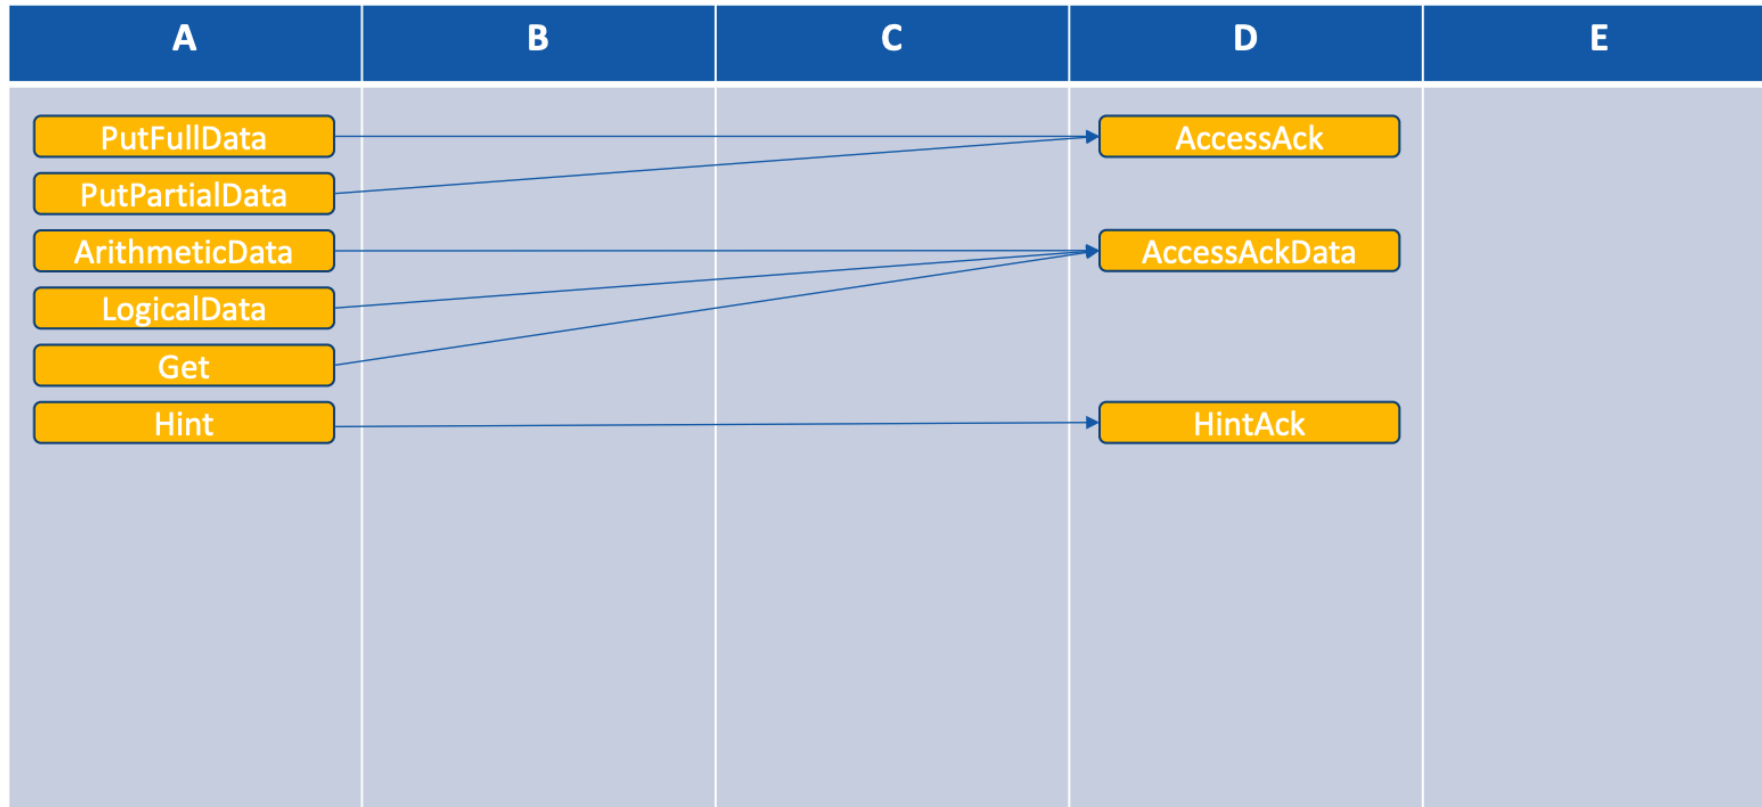
TileLink / TL-UL

- Channel A: memory access request from master to slave
- Channel D: memory access response from slave to master



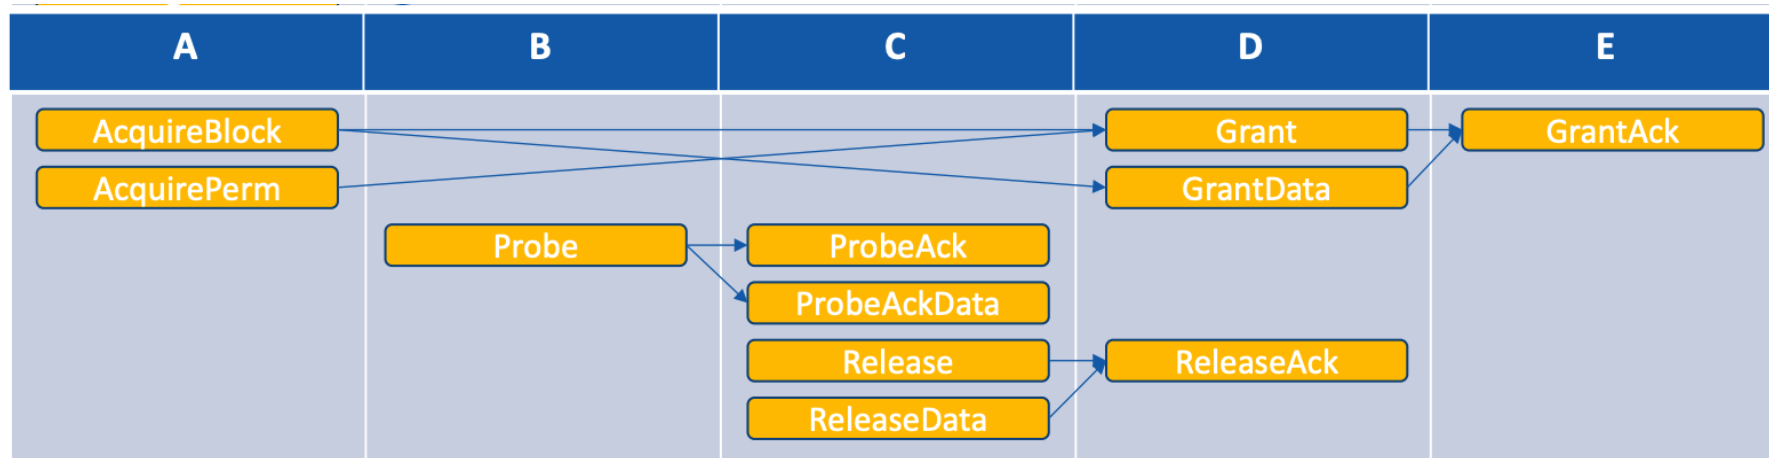
TileLink / TL-UH

- Arithmetic & logical: atomic operation (read-modify-write)
- Hint: prefetch data with intent to read/write



TileLink / TL-C

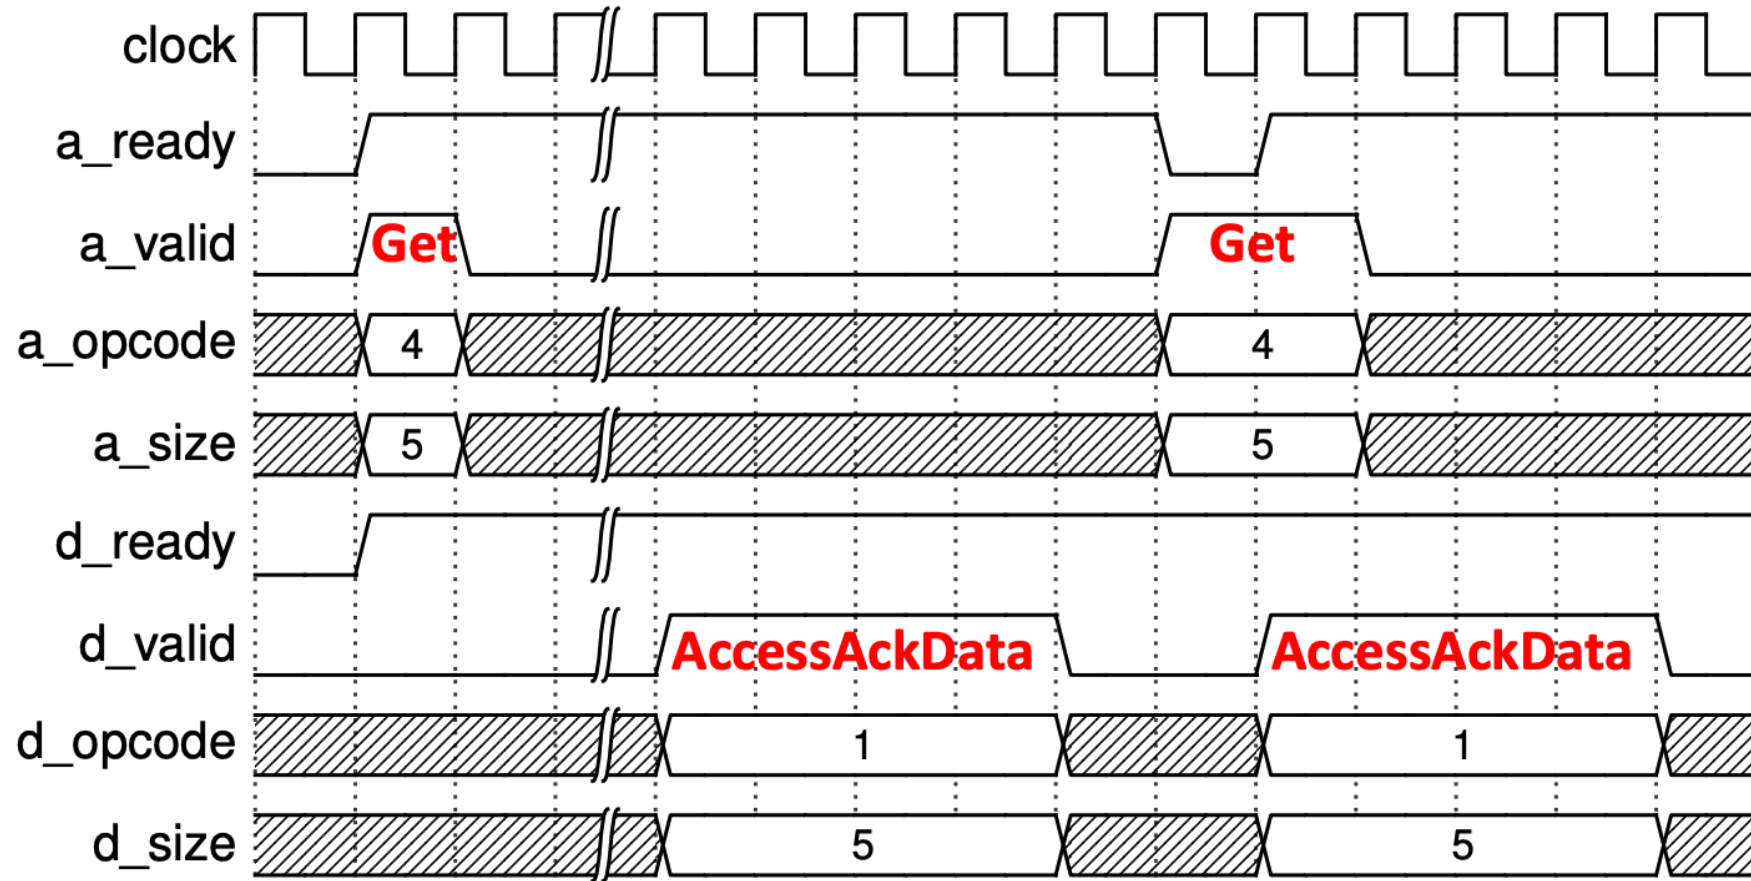
- Acquire: master to slave
 - To obtain cache block to make it local, or write permission of a cache block it already has
- Grand: slave to master
 - As a response to acquire
- Probe: slave to master
 - To query or modify the permission of cache
- Release: master to slave
 - To downgrade its permission of a cache block



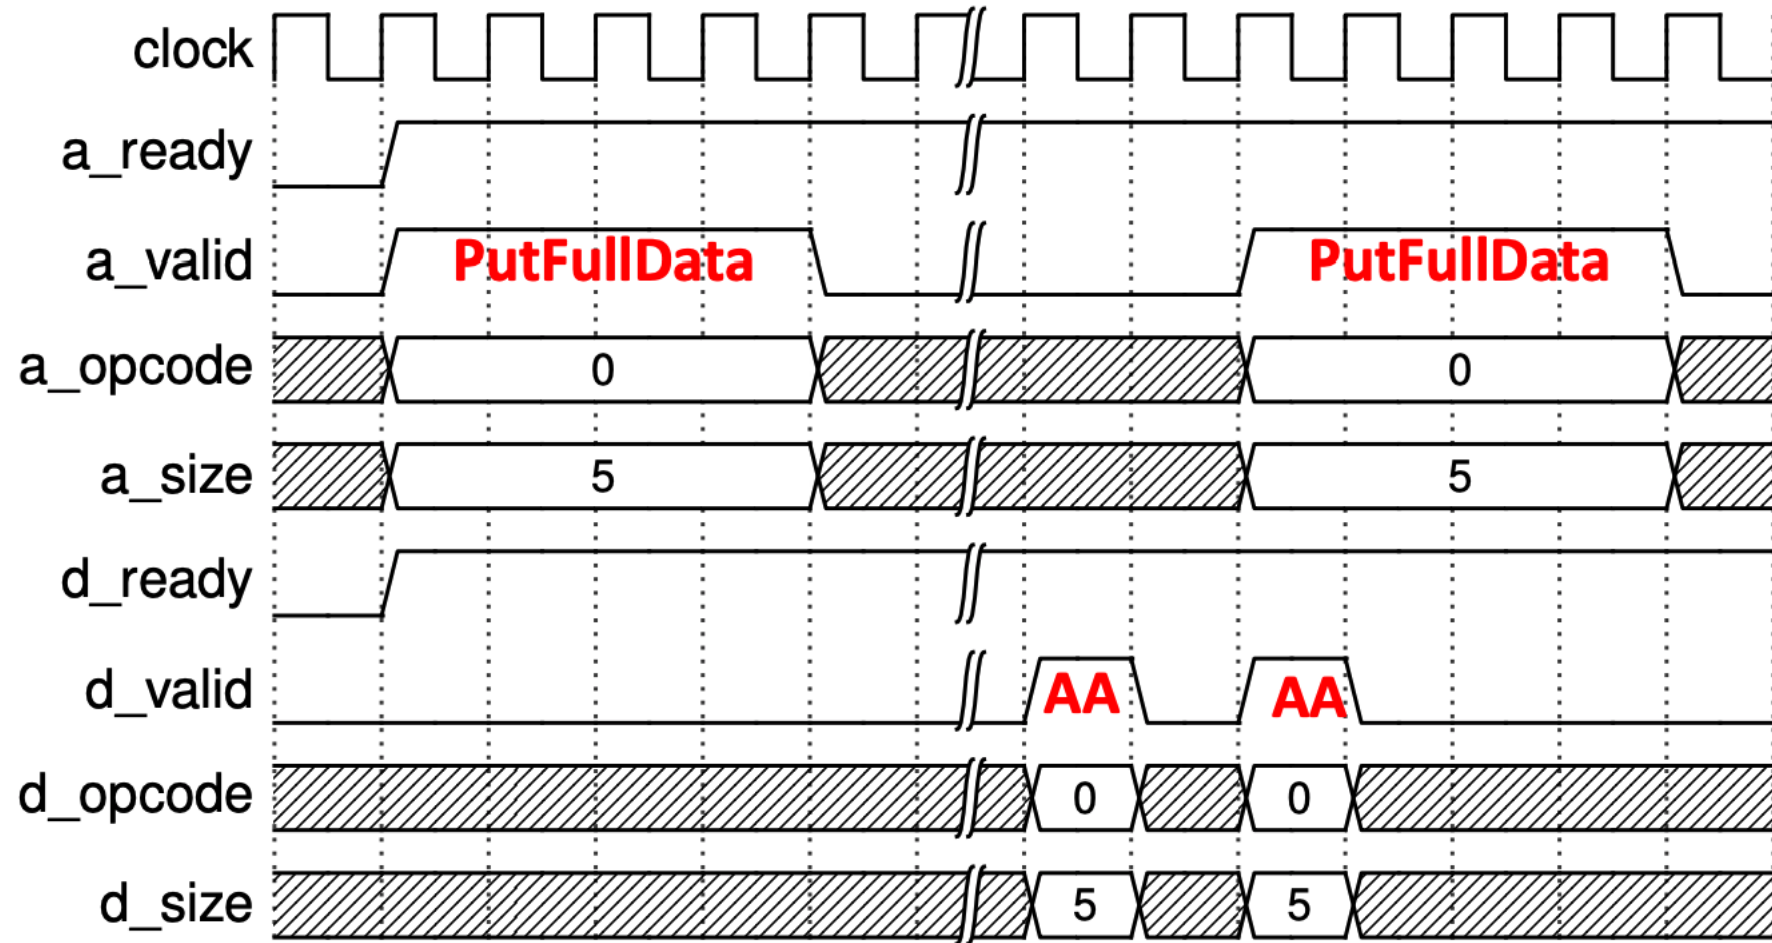
TileLink / wire protocol

- Physically independent channel for each priority (A/B/C/D/E)
- Ready-valid hand-shaking protocol
 - The waiting time could be arbitrary
- Burst transmissions + serialization (variable bus width)
 - Between burst beats, only data field changes
 - Size must align with 2^N , the size field is just the N
 - Number of beats is calculated from size, no special indicator

TileLink / memory read example



TileLink / memory write example



TileLink / cache coherency

Hardware supported cache coherency will be discussed in later session

TileLink / products

SiFive's CPU core complex is built on TileLink. So it's silicon proven.

- Crossbar, adapters, switch
- Snooping-based cache-coherent manager
- Bridges to AMBA protocol

<https://github.com/chipsalliance/rocket-chip/tree/master/src/main/scala/tilelink>

Problem? All in Chisel ...

Table of Content

Uncore

Interrupt Controller

TileLink

>>>> Debug

Summary

Debug / spec

What is debug spec for?

- Debug software, either embedded or kernel
- Help with system bring up before any working CPU on the chip
- Standardize interface between software and hardware debugger

What is debug spec not for?

- Find hardware bugs, but can help to narrow them down

3 parts of the debug system

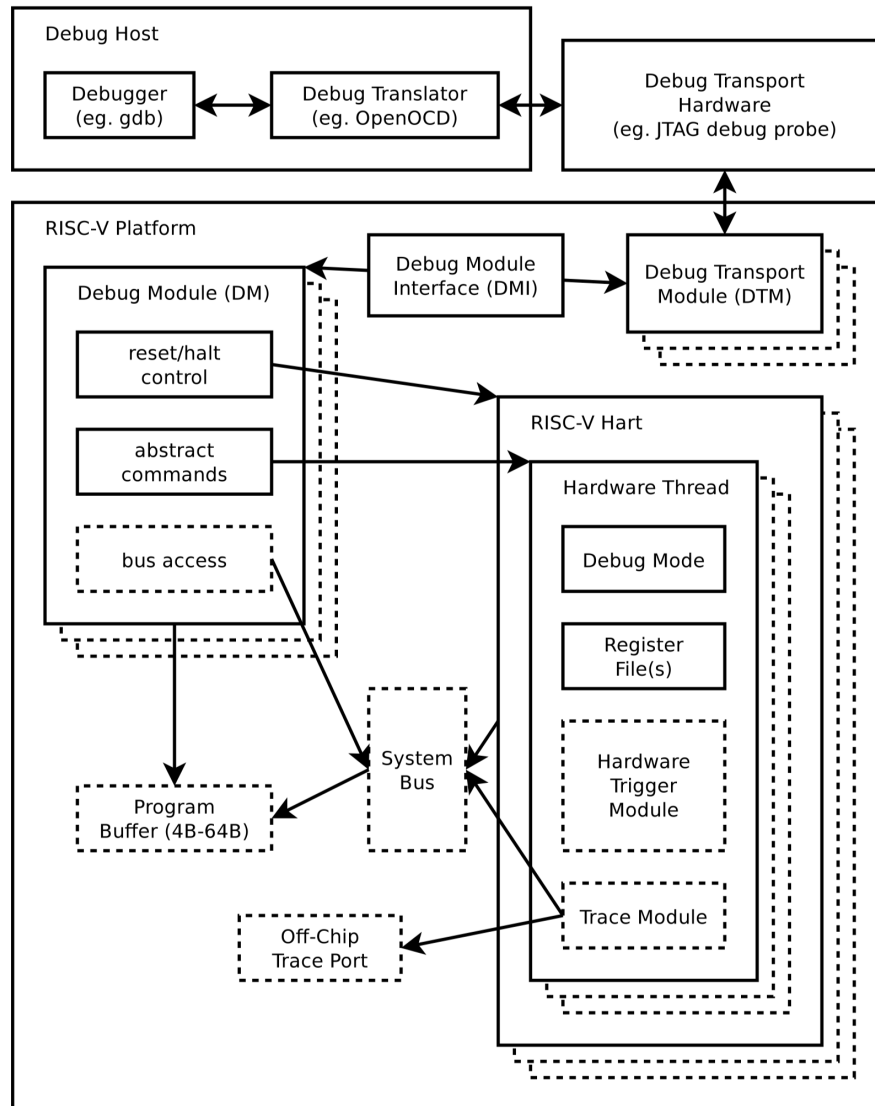
- Debugger software running on the host, like GDB or some GUI IDE
- Debugger hardware connecting host to target, e.g. TRACE32, J-Link
- Silicon block inside the target chip
 - Auto discovery; protocol translation; message passing; runtime control; cross-triggering; etc

Debug spec

principle

- **"Software is King"**
 - Helping software (including kernel) debug is the primary goal
- Decoupled from implementation, both silicon and debugger
- Simple & hardware friendly (again)

Debug spec / system diagram



- Need to distinguish
 - Inside / outside of the chip
 - Software / hardware

Debug / components

DTM (debug transport module)

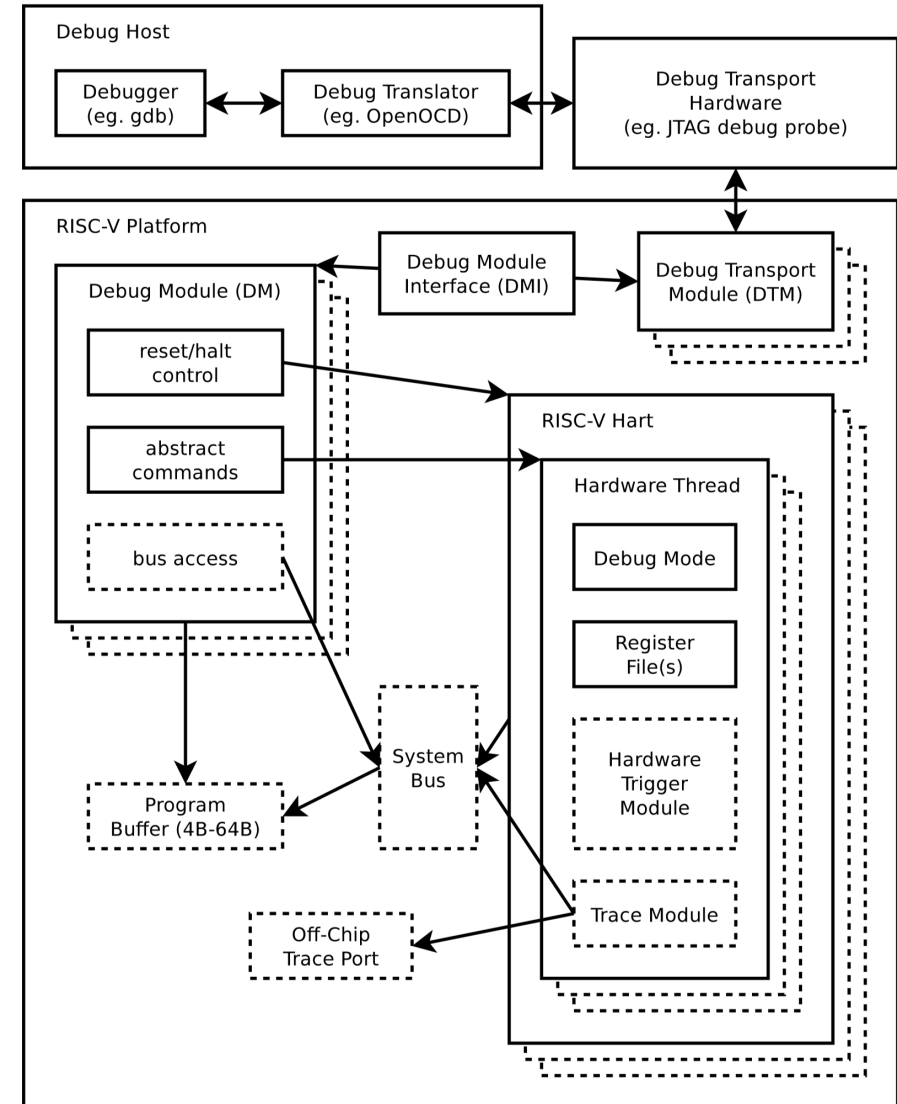
- General interface for different protocols, e.g. JTAG, SPI, USB, etc.
- Right now, only JTAG is defined, but SPI is also supported

DMI (debug module interface)

- Interface bus between DTM and DM
- It is implementation dependent, can be TileLink or AMBA

DM (debug module)

- The core of the debug silicon block
- Provides runtime control, abstract commands, function fabric access



Debug / function of debug module

Reset control

- Output a global reset signal, `ndmreset` (non-debug module reset)
- So that we can debug the first instruction out of reset

Selecting harts

- Debugger can select one or multiple harts to debug

Run control

- Halt, resume, halt-on-reset, reset, singl-step

Abstract commands

- Abstract commands can be injected by debugger to direct CPU carrying out specific commands
- Access registers: GPR, CSR, FPR, etc.
- Quck access: quickly inject some commands for CPU to execute while it's still running
- Access memory

Debug / triggering

How to go into debug mode?

- Breakpoint
 - Also `EBREAK` instruction
- Watchpoint: register
- Cross-trigger
 - Other core encounters a breakpoint
 - Other component hits certain condition

Debug / security

Another big topic in debug is **security**.

- Debug system is powerful, can access almost all the resource
- Big security loophole if left open after production phase

What to do?

- Use fuse bit to disable debug after production phase
- Add authentication step before going into debug mode

Debug / debugger

Software running on the host. No matter has GUI or not.

Open source: OpenOCD + GDB

- Currently only support JTAG interface
- Support runtime control, abstract commands
- Support hardware breakpoints & watchpoints

Commercial

- Lauterbach: TRACE32
- Segger: J-Link
- UltraSoC: UltraDev2

Debug / heterogenous SoC

Discussion: what is in your pocket?

Table of Content

Uncore

Interrupt Controller

TileLink

Debug

>>>> Summary

Summary

Uncore includes blocks that support CPU core's functionality, but not peripheral.

Interrupt controller

- 3 types of interrupts
 - External/software/timer
- PLIC: external interrupt aggregator
- CLINT: software/timer interrupt
- Software/hardware co-workflow

Debug

- Run-control and abstract command
- Decouple implementation of software/hardware, host/target

Cache

- No spec, implementation dependent

Network fabric

- No spec, implementation dependent
- TileLink: not a part of RISC-V ISA spec

Note: ISA includes specs that defines software/hardware interface

감사합니다 Natick
Grazie Danke Ευχαριστίες Dalu
Thank You Köszönöm
Спасибо Dank Gracias
谢谢 Merci Seé
ありがとう Obbrigado